

Generative Models

Unsupervised Learning

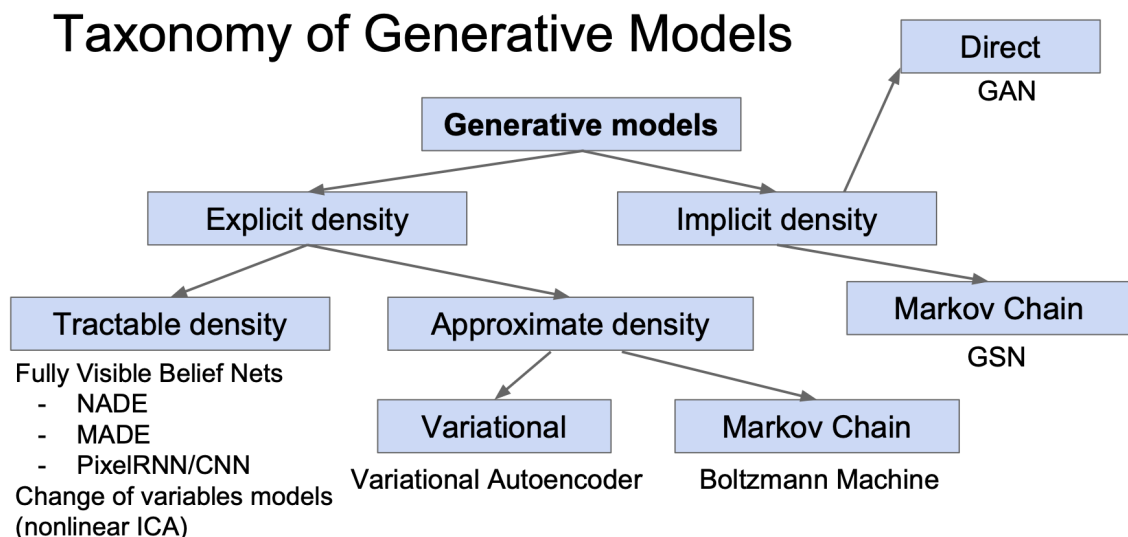
1. Background: Supervised Learning
 - a. Given data (x, y)
 - i. x : data
 - ii. y : label
 - b. Learn a function to map $x \rightarrow y$
 - c. Examples
 - i. classification, regression, object detection, semantic segmentation, image captioning, etc
2. Unsupervised Learning
 - a. Given data x
 - b. Learn some underlying hidden structure of the data
 - c. Examples
 - i. clustering, dimensionality reduction, feature learning, density estimation, etc

Generative Models

1. Generative models
 - a. Given training data
 - b. Generate new samples from same distribution
 - i. want to learn $p_{model}(x)$ similar to $p_{data}(x)$
 - $p_{data}(x) \sim$ training data distribution
 - $p_{model}(x) \sim$ generated sample distribution
2. Density estimation
 - a. 관측된 데이터들의 분포로부터 원래 변수의 확률 분포 특성을 추정하는 것
 - i. 어떤 변수 x 의 밀도 (density)를 추정하는 것은 x 의 확률 밀도 함수 (probability density function)를 추정하는 것과 동일함


- b. For supervised
 - i. 미리 확률 밀도 함수를 정해두고 데이터들로부터 모델의 파라미터만 추정하는 방식 (parametric)
- c. For unsupervised
 - i. 어떠한 사전 정보나 지식 없이 순수하게 관측된 데이터만으로 확률 밀도 함수를 추정하는 방식 (non-parametric)
- d. Density estimation is a core problem in unsupervised learning
 - i. generative model addresses density estimation
- e. Type
 - i. explicit density estimation
 - explicitly define and solve for $p_{model}(x)$
 - ii. implicit density estimation
 - learn model that can sample from $p_{model}(x)$ without explicitly defining it

3. Taxonomy



PixelRNN & PixelCNN

1. PixelRNN
 - a. Generate image pixels starting from corner
 - b. Dependency on previous pixels modeled using an RNN (LSTM)




c. Drawback

- i. sequential generation is slow b/c iteratively go through and add all the pixels

2. PixelCNN

- a. Still generate image pixels starting from corner
- b. Dependency on previous pixels now modeled using a CNN over context region
 - i. train by maximizing the likelihood of training images



c. Advantage

- i. training is faster than PixelRNN b/c can parallelize convolutions since context region values known from training images

d. Drawback

- i. generation must still proceed sequentially, so still slow

3. Pros & Cons

a. Pros

- i. can explicitly compute likelihood $p(x)$
- ii. explicit likelihood of training data gives good evaluation metric
- iii. good samples

b. Cons

- i. sequential generation is too slow

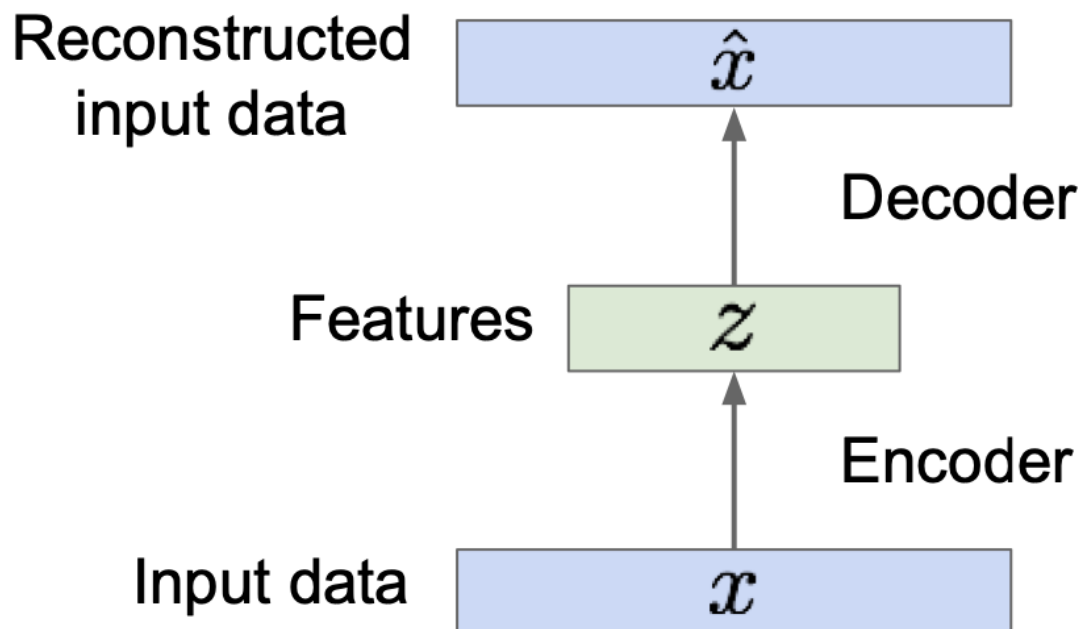
Variational Autoencoders (VAE)

1. PixelCNN vs. VAE

- a. PixelCNNs define tractable density function, optimize likelihood of training data $p_{\theta}(x)$
- b. VAEs define intractable density function with latent z
 - i. cannot optimize directly, derive and optimize lower bound on likelihood, instead

2. Background: Autoencoder

- a. Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data



b. Encoder

- i. generate z from x
 - z usually smaller than x b/c dimensionality reduction
 - dimensionality reduction 을 통해 중요한 feature만 z 로 선정됨

c. Decoder

- i. reconstruct \hat{x} from z

d. Process

- i. Train using L2 loss function

- $\|x - \hat{x}\|^2$

ii. Throw away decoder

iii. Use encoder to initialize a supervised model

- softmax loss function and fine-tuning

e. Purpose

- Autoencoder의 decoder 부분은 encoder를 학습시키는 것이 목적
- 데이터 생성이 목적이 아님
- 데이터를 잘 압축하는 것 (i.e., 데이터의 특징을 잘 추출하는 것, 데이터의 차원을 잘 줄이는 것) 이 목적

3. Variational autoencoder

a. Purpose

- VAE는 decoder를 학습시키기 위해 encoder를 사용함
- decoder를 이용하여 데이터를 생성하는 것이 목적

b. Assume training data $\{x^{(i)}\}_{i=1}^N$ is generated from underlying latent representation z

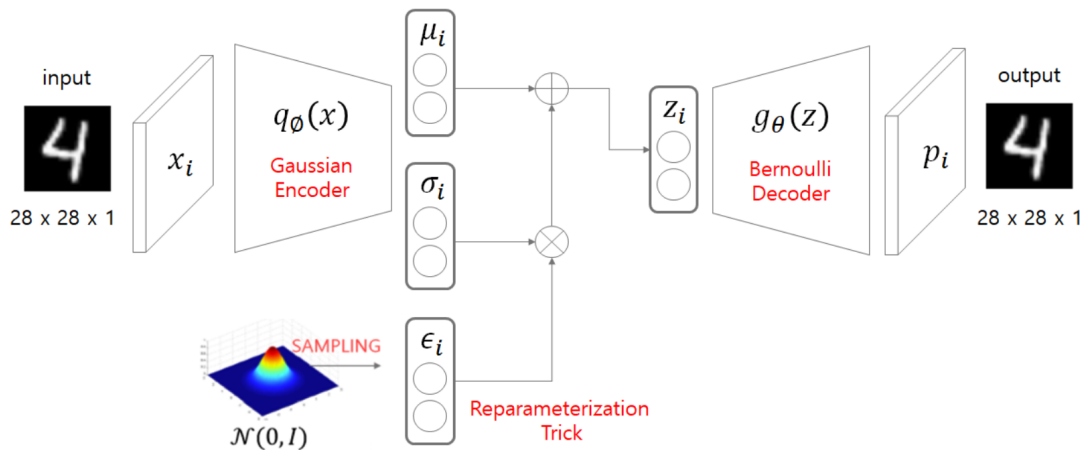
- z is sampled from true prior $p_{\theta^*}(z)$
- x is sampled from true conditional $p_{\theta^*}(x|z^{(i)})$

c. Estimate the true parameters θ^* of the generative model

d. How?

- choose prior $p(z)$ to be simple, e.g., Gaussian
- conditional $p(x|z)$ is complex (generates image), so present with neural network

e. Overall structure



f. Encoder

- i. 데이터 x 를 입력으로 넣어서 데이터의 특징을 추측함
- ii. 이 때, 특징들의 분포는 정규 분포를 따른다고 가정함
- iii. 따라서 아웃풋으로 특징 분포의 μ_i, σ_i 를 리턴함
- iv. 문제점: 분포 내 데이터를 어떻게 뽑아서 사용할까?
 - reparameterization trick을 사용하여 샘플링함
 - 직접 샘플링: 가우시안 분포에서 직접 샘플링을 하면 역전파가 불가능함
 - reparameterization trick: ϵ 을 $\mathcal{N}(0, 1)$ 에서 샘플링해서 σ_i 와 곱하고 μ_i 와 더하여 같은 확률 분포에서 값을 샘플링함

g. Decoder

- i. 특징 z 를 통해 데이터 x 를 복원함 (z 를 줬을 때 x 가 나올 확률을 maximize 해야함)
- ii. loss function
 - reconstruction error: 복원된 샘플이 얼마나 원본과 유사하도록 만들어줌
 - regularization term: 이상적인 z 의 분포와 encoder를 통해 얻은 z 의 분포가 유사하도록 만들어줌 (KL-divergence)

4. Pros & Cons

a. Pros

- i. principled approach to generative models

- ii. allows inference of $q(z|x)$, can be useful feature representation for other tasks
- b. Cons
 - i. not as good evaluation as PixelRNN or PixelCNN
 - ii. samples blurrier and lower quality compared to SOTA (GANs)

Generative Adversarial Networks (GAN)

1. Overview

- a. GANs don't work with any explicit density function
- b. Take game-theoretic approach, instead
 - i. learn to generate from training distribution through 2-player game

2. GANs

a. Problem

- i. want to sample from complex, high-dimensional training distribution

b. Solution

- i. sample from a simple distribution, e.g., random noise

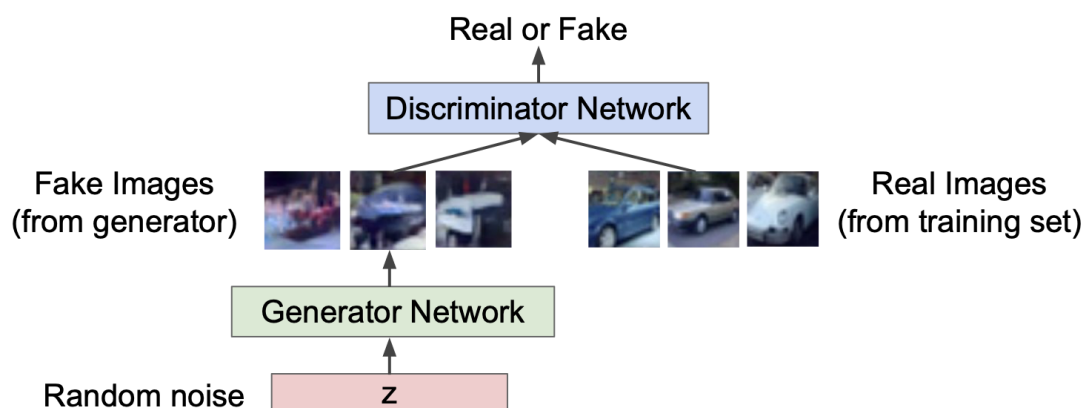
c. Training

i. Generator network

- try to fool the discriminator by generating real-looking images

ii. Discriminator network

- try to distinguish between real and fake images



- d. After training
 - i. Use generator network to generate new images
- 3. Pros & Cons
 - a. Pros
 - i. SOTA samples
 - b. Cons
 - i. more unstable to train
 - ii. can't solve inference queries such as $p(x)$, $p(z|x)$